

RAG for Agentic AI

★ What is RAG (Retrieval-Augmented Generation)?

RAG — *Retrieval-Augmented Generation* — is an AI technique where a Large Language Model (LLM) combines **external knowledge** with its own reasoning to produce accurate, reliable answers.

LLM + Your Documents = RAG

RAG fixes the biggest limitations of LLMs:

- They **don't know** your PDF, website, or database.
- They **forget** facts after 2023 cutoff (in many models).
- They **hallucinate** when unsure.

RAG solves this by providing **real, verified information** from your documents at the time of answering.

★ Why RAG is Important?

RAG makes AI:

- ✓ More accurate
- ✓ Less hallucinating
- ✓ More trustworthy
- ✓ Capable of using your custom data
- ✓ Scalable for enterprise knowledge systems

✿ PART 1 — DOCUMENTS (Knowledge Source)

Short Theory (Simple Enough for Notes)

Documents = any knowledge your RAG system uses to answer questions.

Examples:

- PDFs
- Markdown, .txt
- Websites
- Databases
- CSV, JSON

- API responses

LLMs **cannot** store your private/company knowledge → so RAG uses **your documents** as the source of truth.

Before RAG can use a document, it must be:

- 1 Loaded
- 2 Cleaned
- 3 Structured
- 4 Chunked
- 5 Embedded
- 6 Stored in vector DB

Code: Loading Documents (PDF, Text, Website)

1. Load PDF (most common)

```
from langchain.document_loaders import PyPDFLoader

loader = PyPDFLoader("policy.pdf")
documents = loader.load()

print(documents[0].page_content[:300])
```

2. Load Text / Markdown

```
from langchain.document_loaders import TextLoader

loader = TextLoader("notes.txt")
documents = loader.load()
```

3. Load Website Page

```
from langchain.document_loaders import WebBaseLoader

loader = WebBaseLoader("https://example.com/docs")
documents = loader.load()
```

4. Load Data from CSV / DB Rows

```
import pandas as pd

from langchain.schema import Document

df = pd.read_csv("data.csv")

documents = [
    Document(page_content=row["text"], metadata={"id": i})
    for i, row in df.iterrows()
]
```

5. Clean the Document (remove junk)

Basic cleaning improves RAG accuracy.

```
def clean(text):

    text = text.replace("\n", " ")
    text = " ".join(text.split())

    return text

cleaned_docs = []

for d in documents:
    d.page_content = clean(d.page_content)
    cleaned_docs.append(d)
```

💡 Metadata (very important for Agentic RAG)

Metadata helps agents' reason better.

```
for i, doc in enumerate(cleaned_docs):
    doc.metadata["chunk_id"] = i
    doc.metadata["source"] = "policy.pdf"
```

PART 2 — CHUNKING

What is Chunking? (Short Theory)

Chunking = breaking large documents into **small readable pieces** (chunks) so that:

- LLM can understand the text
- embeddings become more meaningful
- retrieval becomes accurate
- agent can reason step-by-step

Ideal Chunk Size:

- **350–800 tokens** (most used)
- With **10–20% overlap**

Why overlap?

To prevent loss of meaning between paragraphs.

Why Chunking Is Critical for Agentic AI?

Agents need context that is:

- ✓ Accurate
- ✓ Relevant
- ✓ Reusable
- ✓ Consistent

Good chunking = Good retrieval → Better decisions by agents.

CODE: Simple Fixed-Size Chunking

```
from langchain.text_splitter import CharacterTextSplitter

splitter = CharacterTextSplitter(
    chunk_size=500,
    chunk_overlap=100,
    separator="\n"
)

chunks = splitter.split_documents(cleaned_docs)
print("Total Chunks:", len(chunks))
```

This is the **best & most stable** chunking method.

```
from langchain.text_splitter import RecursiveCharacterTextSplitter

splitter = RecursiveCharacterTextSplitter(
    chunk_size=700,
    chunk_overlap=100,
    separators=["\n\n", "\n", " ", ""]
)

chunks = splitter.split_documents(cleaned_docs)
```

This helps agents understand where the chunk came from.

```
for i, chunk in enumerate(chunks):
    chunk.metadata["chunk_id"] = i
    chunk.metadata["length"] = len(chunk.page_content)
```

🔍 Chunk Quality Improvement Tricks

✓ 1. Use Sentence-Aware Chunking (optional)

Good when documents have long paragraphs.

```
from langchain.text_splitter import NLTKTextSplitter

splitter = NLTKTextSplitter(chunk_size=600)
chunks = splitter.split_documents(cleaned_docs)
```

✓ 2. Clean Junk Words Before Chunking

Otherwise junk spreads into every chunk's embedding.

✓ 3. Remove Small Chunks (<200 chars)

They reduce retrieval accuracy.

```
chunks = [c for c in chunks if len(c.page_content) > 200]
```

Example

Chunk ID: 42

Source: policy.pdf

Text: "Employees are allowed 24 paid leaves per year..."

Length: 532

PART 3: Retrieval Engine

This part explains **how your agent finds the right chunk** when you ask a question.

RAG Retrieval =

User Query → Embedding → Vector DB Search → Relevant Chunks

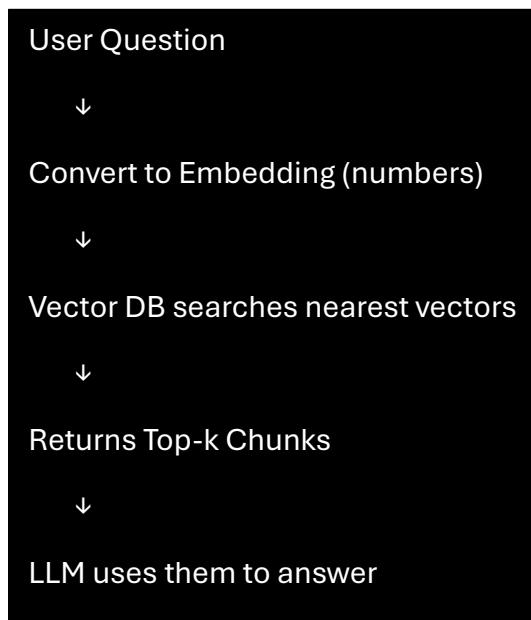
1. What is Retrieval?

Retrieval = “The process of pulling correct knowledge from your DB”.

It has 2 parts:

1. **Vector Search (ANN - Approx Nearest Neighbor)**
2. **Re-ranking** (optional but improves accuracy)

2. Retrieval Flow



3. Retrieval Parameters You Must Know

Parameter	Meaning	Recommended
k	number of chunks to retrieve	3–5
similarity metric	cosine / euclidean / dot	cosine
score_threshold	ignore low-quality matches	0.2–0.3
reranking	reorder best chunks using LLM	optional but powerful

4. Types of Retrieval

1 Dense Vector Search (FAISS / Chroma)

- Best for most RAG systems
- Uses embeddings
- Simple & fast

2 Sparse Retrieval (BM25)

- Word-based search
- Great for code + exact matching
- Often used in hybrid RAG

3 Hybrid Retrieval

Embedding + BM25 combined

 *Best accuracy, used in advanced Agentic AI*

★ PART 3 — CODE SECTION

✓ Code #1 — Vector Search with FAISS

```
from langchain_community.vectorstores import FAISS

from langchain_openai import OpenAIEMBEDDINGS

emb = OpenAIEMBEDDINGS(model="text-embedding-3-small")

# Create vector store

vectorstore = FAISS.from_documents(chunks, embedding=emb)

# Save locally

vectorstore.save_local("faiss_rag_db")
```

✓ Code #2 — Search Relevant Chunks

```
query = "What is RAG and how does retrieval work?"

results = vectorstore.similarity_search(
    query=query,
    k=3
)

for r in results:
    print("\n---Chunk---\n", r.page_content)
```

✓ **Code #3 — Search With Score Threshold**

```
docs = vectorstore.similarity_search_with_score(  
    query="Explain agentic AI",  
    k=5  
)  
  
filtered = [d for d, score in docs if score < 0.25]  
  
print("Returned:", len(filtered))
```

★ **5. Add Re-ranking (Boost Accuracy)**

Your LLM can reorder results based on relevance.

✓ **Code #4 — Simple Reranking with LLM**

```
from langchain_openai import ChatOpenAI  
  
llm = ChatOpenAI(model="gpt-4o-mini")  
  
  
def rerank(query, retrieved_chunks):  
  
    prompt = f"""\n        Query: {query}\n        Rank the following chunks by relevance:\n        {retrieved_chunks}\n        Return top 3.\n        """  
  
    return llm.invoke(prompt)  
  
  
reranked = rerank(query, [c.page_content for c in results])  
print(reranked)
```

⭐ 6. Hybrid Retrieval (Best for Agentic AI)

Dense + Sparse combined → more accurate.

✓ Code #5 — Hybrid Search

```
from langchain_community.retrievers import BM25Retriever  
from langchain.retrievers import EnsembleRetriever  
dense_retriever = vectorstore.as_retriever(search_kwargs={"k": 3})  
sparse_retriever = BM25Retriever.from_documents(documents)  
hybrid = EnsembleRetriever(  
    retrievers=[dense_retriever, sparse_retriever],  
    weights=[0.6, 0.4]  
)  
out = hybrid.invoke("how does embedding work?")
```

📌 Retrieval Best Practices

- Use FAISS for local RAG experiments
- Use Chroma for production
- Always retrieve top-k=3 to avoid hallucination

🔥 RAG Skill Tree – PART 4

“LLM Integration + Answer Generation Pipeline”

⭐ PART 4: LLM + Generation Engine

This is where RAG becomes *useful*.

You already learned:

- Load Docs
- Chunk Docs
- Create Embeddings

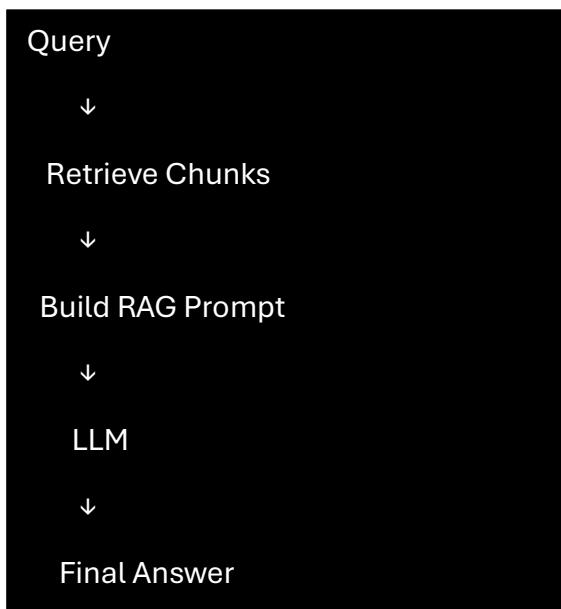
- Build Vector DB
- Retrieve Relevant Chunks

Now → we plug it into the **LLM** to generate accurate answers.

RAG =

Retrieve (External Knowledge) + Generate (LLM reasoning)

💡 1. How Answer Generation Works (Simple Visual)



💡 2. What is a RAG Prompt?

It is a special prompt where we give the LLM:

- ✓ user query
- ✓ retrieved chunks
- ✓ instructions
- ✓ restrictions

Example:

Use ONLY the provided context to answer.
If answer not found, say "Not in document".

Context:

[chunk1]
[chunk2]
[chunk3]

User question:

Explain RAG retrieval.

This eliminates hallucination.

This turns the LLM into a **grounded agent**.

💡 3. Code: Build a Simple RAG Pipeline

This is the **standard RAG code** used everywhere.

✓ Code #1 — Setup LLM

```
from langchain_openai import ChatOpenAI  
llm = ChatOpenAI(model="gpt-4o-mini")
```

✓ **Code #2 — Build RAG Chain**

```
def build_rag_prompt(query, retrieved_docs):  
  
    context = "\n\n".join([d.page_content for d in retrieved_docs])  
  
    prompt = f"""\n  
    Use ONLY the context below to answer the question.  
  
    Context:  
  
    {context}  
  
    Question:  
  
    {query}  
  
    If not found in context, reply: 'Information not available.'  
    """  
  
    return prompt
```

✓ **Code #3 — RAG Execution**

```
query = "What is agentic AI?"  
  
# Step 1: Retrieve  
docs = vectorstore.similarity_search(query, k=3)  
  
# Step 2: Build prompt  
rag_prompt = build_rag_prompt(query, docs)
```

This is a **complete RAG answer generator**.

💡 4. Optional but Powerful: RAG with Chain

LangChain gives a ready-built pipeline.

✓ Code #4 — LangChain RAG Chain

```
from langchain.chains import RetrievalQA

rag_chain = RetrievalQA.from_chain_type(
    llm=llm,
    retriever=vectorstore.as_retriever(search_kwargs={"k": 3}),
    return_source_documents=True
)
The out = rag_chain.invoke({"query": "Explain embeddings"})
print(out["result"])
```

- builds prompt
 - merges text
 - produces answer
-

💡 5. Improve Output with Instructions (Most Important)

Your LLM produces better results with strong guiding instructions.

RAG System Instructions Template

```
You are a highly accurate retrieval-focused AI.
```

Rules:

- Use ONLY the provided context.
- Do NOT add external knowledge.
- Cite the chunk if necessary.
- If answer not found, say "Not found in the document."
- Keep the answer clear and concise.

6. Add Safety: Prevent Hallucination

Add this rule:

If context does not contain relevant info:

Return: "Not available in the provided document."

This makes your RAG **robust** for Agentic AI.

7. Outputs You Can Generate

RAG can produce:

- Summary
- Q&A
- Explanations
- Step-by-step instructions
- Convert PDF → Structured data
- Reasoning based on retrieved chunks

8. Full RAG Pipeline

Full RAG Answer Generation Flow

- 1 User asks a question
- 2 Convert query → embedding
- 3 Vector DB performs semantic search
- 4 Top-k chunks returned
- 5 Build final RAG prompt
- 6 LLM generates grounded answer
- 7 Optional: citations + re-ranking

9. Full Working RAG Script (Simple)

```
from langchain_community.vectorstores import FAISS
from langchain_openai import OpenAIEmbeddings, ChatOpenAI
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_community.document_loaders import PyPDFLoader

# 1 Load
loader = PyPDFLoader("myfile.pdf")
docs = loader.load()

# 2 Chunk
splitter = RecursiveCharacterTextSplitter(chunk_size=400, chunk_overlap=50)
chunks = splitter.split_documents(docs)

# 3 Embed
emb = OpenAIEmbeddings(model="text-embedding-3-small")

# 4 Vector DB
vectorstore = FAISS.from_documents(chunks, emb)

# 5 LLM
llm = ChatOpenAI(model="gpt-4o-mini")

# 6 RAG Pipeline
def rag(query):
    docs = vectorstore.similarity_search(query, k=3)
    context = "\n\n".join([d.page_content for d in docs])
```

```
prompt = f"""
Answer using ONLY this context.

If not found, say 'Not available.'
```

Context:

```
{context}
```

Question: {query}

"""

```
return llm.invoke(prompt).content
```

```
print(rag("Explain RAG retrieval"))
```

⌚ Part 4 — What You Learned

LLM + Retrieval = RAG Engine

- ✓ How generation works
 - ✓ Prompt engineering for RAG
 - ✓ Build full RAG chain
 - ✓ Safety rules
 - ✓ Complete working code
 - ✓ Perfect for Agentic AI answering systems
-